

```
[4]: import findspark
findspark.init()
```

```
[5]: from pyspark.sql import SparkSession
from pyspark.sql import SQLContext
spark = SparkSession.builder.master("local").appName("Test Spark").getOrCreate()
sc=spark.sparkContext
sql=SQLContext(sc)
```

```
C:\BIG DATA\spark\spark1\spark2\python\pyspark\sql\context.py:112: FutureWarning: Deprecated in 3.0.0. Use SparkSession.builder.getOrCreate() instead.
warnings.warn(
```

```
[6]: spark
```

```
[6]: SparkSession - in-memory
```

SparkContext

[Spark UI](#)

Version	v3.3.1
Master	local
AppName	Test Spark

```
[17]: import pandas as pd
```

```
[93]: train_dataset, test_dataset = finalized_data.randomSplit([0.7, 0.3])
      from pyspark.ml.regression import LinearRegression
      MLR = LinearRegression(featuresCol="features", labelCol="Profit")
      model = MLR.fit(train_dataset)
      pred = model.evaluate(test_dataset)
      pred.predictions.show()
```

```
+-----+-----+-----+
|          features|    Profit|    prediction|
+-----+-----+-----+
|[0.0,0.0,115949.7...| 134448.9|140303.95433473212|
|[0.0,0.0,117590.2...|  59672.75| 67958.53017332428|
|[0.0,1.0,53057.14...|100131.14|101420.10718197195|
|[0.0,1.0,97963.63...|  81680.49| 80917.64640287402|
|[0.0,1.0,123371.5...|137174.93|125486.08475656048|
|[0.0,1.0,137269.0...|144489.35|135453.29285333218|
|[1.0,0.0,53517.15...|  45855.41| 63724.53964451117|
|[1.0,0.0,129156.3...|  79940.98| 71499.63916226041|
|[1.0,0.0,154475.9...|107665.56|107554.83584866446|
+-----+-----+-----+
```

```
[94]: coefficient = model.coefficients
      print ("The coefficients of the model are : %a" %coefficient)
      intercept = model.intercept
      print ("The Intercept of the model is : %f" %intercept)
```

```
The coefficients of the model are : DenseVector([-738.216, -2607.6941, -0.0666, 0.7994, 0.0241])
The Intercept of the model is : 65779.728854
```

```
[95]: from pyspark.ml.evaluation import RegressionEvaluator
      evaluation = RegressionEvaluator(labelCol="Profit", predictionCol="prediction")

      r2 = evaluation.evaluate(pred.predictions, {evaluation.metricName: "r2"})
      print("The value of r2-coefficient of determination is : %.3f" %r2)
```

```
[28]: #Setting up the simple linear regression algorithm
from pyspark.ml.regression import LinearRegression
MLALGO=LinearRegression(featuresCol="Features",labelCol="Grades")
model=MLALGO.fit(train_data)
predict=model.evaluate(test_data)
predict.predictions.show()
```

```
+-----+-----+-----+
|Features|Grades|prediction|
+-----+-----+-----+
| [1.0] | 1.5 | 1.552976190476186 |
| [2.0] | 1.8 | 1.8262235449735411 |
| [2.0] | 1.8 | 1.8262235449735411 |
| [4.0] | 2.4 | 2.3727182539682508 |
| [4.0] | 2.4 | 2.3727182539682508 |
| [5.0] | 2.7 | 2.6459656084656062 |
| [6.0] | 2.9 | 2.919212962962961 |
| [6.0] | 2.9 | 2.919212962962961 |
| [7.0] | 3.1 | 3.192460317460316 |
| [9.0] | 3.9 | 3.738955026455026 |
| [11.0] | 4.3 | 4.2854497354497365 |
| [14.0] | 5.0 | 5.105191798941801 |
| [14.0] | 5.0 | 5.105191798941801 |
+-----+-----+-----+
```

```
[30]: #finding coefficient/gradient value
coefficient=model.coefficients
print(" The coefficient value of the dataset is : %a" %coefficient)
```

```
The coefficient value of the dataset is : DenseVector([0.2732])
```

```
[30]: #finding coefficient/gradient value
coefficient=model.coefficients
print(" The coefficient value of the dataset is : %a" %coefficient)
```

The coefficient value of the dataset is : DenseVector([0.2732])

```
[32]: #finding the y intercept value
intercept=model.intercept
print(" The y intercept value of the dataset is: %f" %intercept)
```

The y intercept value of the dataset is: 1.279729

```
[37]: # printing the root mean square error
from pyspark.ml.evaluation import RegressionEvaluator
evaluation=RegressionEvaluator(labelCol="Grades",predictionCol="prediction")
rmse=evaluation.evaluate(predict.predictions, {evaluation.metricName: "rmse"})
print("the root mean square error is :%.3f" %rmse)
```

the root mean square error is :0.071

```
[38]: #printing the mean square error
mse=evaluation.evaluate(predict.predictions,{evaluation.metricName: "mse"})
print(" the mean square error is : %.3f "%mse)
```

the mean square error is : 0.005

```
[39]: mae=evaluation.evaluate(predict.predictions,{evaluation.metricName: "mae"})
print(" the mean absolute error is : %.3f "%mae)
```

the mean absolute error is : 0.056

```
[40]: r2= evaluation.evaluate(predict.predictions,{evaluation.metricName: "r2"})
print(" the r2 coefficient is : %.3f "%r2)
```

the r2 coefficient is : 0.996

```
[77]: from pyspark.ml.classification import LogisticRegression
log_reg=LogisticRegression(labelCol='HeartDisease').fit(training_df)
train_results=log_reg.evaluate(training_df).predictions
train_results.show()
```

features	HeartDisease	rawPrediction	probability	prediction
(9,[0,1,2,3],[40....	1	[-1.8370295339444...	[0.13740298326385...	1.0
(9,[0,1,2,3],[48....	1	[-2.8094399362242...	[0.05681618611795...	1.0
(9,[0,1,2,3],[48....	1	[-2.5808761497936...	[0.07037938622252...	1.0
(9,[0,1,2,3],[51....	1	[-3.1592249137334...	[0.04072932576133...	1.0
(9,[0,1,2,3],[51....	1	[-3.0143754950465...	[0.04678064527891...	1.0
(9,[0,1,2,3],[53....	1	[-2.8606616500488...	[0.05413281238100...	1.0
(9,[0,1,2,3],[59....	1	[-3.3899746998493...	[0.03261025343674...	1.0
(9,[0,1,3],[38.0,...	1	[-0.1615377473016...	[0.45970315207475...	1.0
(9,[0,1,3],[39.0,...	1	[-0.4477509315172...	[0.38989563510280...	1.0
(9,[0,1,3],[41.0,...	1	[-0.2394819145126...	[0.44041402870399...	1.0
(9,[0,1,3],[42.0,...	1	[-1.3059459798178...	[0.21316601909992...	1.0
(9,[0,1,3],[44.0,...	1	[-0.8606401774520...	[0.29720561159364...	1.0
(9,[0,1,3],[46.0,...	1	[-1.0431863780333...	[0.26053564820238...	1.0
(9,[0,1,3],[46.0,...	1	[-1.1696987345155...	[0.23690944374255...	1.0
(9,[0,1,3],[46.0,...	1	[-1.2294312108545...	[0.22628099263843...	1.0
(9,[0,1,3],[49.0,...	1	[-1.5406489439600...	[0.17644095722630...	1.0
(9,[0,1,3],[49.0,...	1	[-0.5351414812995...	[0.36931851958282...	1.0
(9,[0,1,3],[50.0,...	1	[-0.8683949088494...	[0.29558839780709...	1.0
(9,[0,1,3],[51.0,...	1	[-2.1187316518522...	[0.10728949041858...	1.0
(9,[0,1,3],[52.0,...	1	[-0.8858615085448...	[0.29196460451822...	1.0

only showing top 20 rows

```
[78]: results=log_reg.evaluate(test_df).predictions
results.show()
```

```
[76]: categorical_cols = [item[0] for item in datasets.dtypes if item[1].startswith('string')]
print(categorical_cols)
numerical_cols = [item[0] for item in datasets.dtypes if item[1].startswith('int') | item[1].startswith('double')][:-1]
print(numerical_cols)
print(str(len(categorical_cols)) + ' categorical features')
print(str(len(numerical_cols)) + ' numerical features')

['City']
['Miscellaneous_Expenses', 'Food_Innovation_Spend', 'Advertising']
1 categorical features
3 numerical features
```

```
[87]: from pyspark.ml.feature import StringIndexer, OneHotEncoder, VectorAssembler
stages = []
for categoricalCol in categorical_cols:
    stringIndexer = StringIndexer(inputCol = categoricalCol, outputCol = categoricalCol + 'Index')
    OHencoder = OneHotEncoder(inputCols=[stringIndexer.getOutputCol()], outputCols=[categoricalCol + "_catVec"])
    stages += [stringIndexer, OHencoder]
assemblerInputs = [c + "_catVec" for c in categorical_cols] + numerical_cols
Vectassembler = VectorAssembler(inputCols=assemblerInputs, outputCol="features")
stages += [Vectassembler]

from pyspark.ml import Pipeline
cols = datasets.columns
pipeline = Pipeline(stages = stages)
pipelineModel = pipeline.fit(datasets)
datasets = pipelineModel.transform(datasets)
selectedCols = ['features'] + cols
datasets = datasets.select(selectedCols)
pd.DataFrame(datasets.take(5), columns=datasets.columns)
```

```
[80]: # true positive
      tp = results[(results.HeartDisease == 1) & (results.prediction == 1)].count()
      tp
```

[80]: 106

```
[81]: #true negative
      tn = results[(results.HeartDisease == 0) & (results.prediction == 0)].count()
      tn
```



[81]: 80

```
[82]: #false positive
      fp = results[(results.HeartDisease == 0) & (results.prediction == 1)].count()
      fp
```

[82]: 23

```
[83]: #false negative
      fn = results[(results.HeartDisease == 1) & (results.prediction == 0)].count()
      fn
```

[83]: 17

```
[84]: #accuracy
      accuracy=float((tp+tn)/(results.count()))
      accuracy
```

[84]: 0.8230088495575221

Ambari

Dashboard

Services

HDFS

YARN

MapReduce2

Tez

Hive

HBase

Pig

Sqoop

Oozie

ZooKeeper

Storm

Infra Solr

Atlas

Kafka

Knox

Ranger

Dashboard / Metrics

Sandbox

admin

METRICSHEATMAPSCONFIG HISTORY

METRIC ACTIONS

LAST 1 HOUR

NameNode Heap

10%

HDFS Disk Usage

69%

NameNode CPU WIO

n/a

DataNodes Live

1/1

NameNode RPC

0.80 ms

Memory Usage

No Data Available

Network Usage

No Data Available

CPU Usage

No Data Available

Cluster Load

No Data Available

NameNode Uptime

1d 3h 20m

HBase Master Heap

n/a

HBase Ave Load

n/a

Region In Transition

HBase Master Uptime

ResourceManager Heap

NodeManagers Live

To quickly get started with the Hortonworks Sandbox, follow this tutorial <https://hortonworks.com/tutorial/hadoop-tutorial-getting-started-with->

To initiate your Hortonworks Sandbox session, open a browser to this address:

For VirtualBox:


Welcome screen: <http://localhost:1080>

```
SSH: http://localhost:4200
```

For VMware:

```
Welcome screen:  http://192.168.100.91:1080
```

```
SSH: http://192.168.100.91:4200
```

 Ambari

Dashboard

Services

HDFS

YARN

MapReduce2

Tez

Hive

HBase

Pig

Sqoop

Oozie

ZooKeeper
















Storm

Infra Solr

Atlas

Kafka

Knox

Name >	Size >	Last Modified >	Owner >	Group >	Permission	Erasure Coding	Encrypted
							
 T2021_2.csv	94.9 MB	2023-01-18 23:20	admin	hdfs	-rw-rw-rw-		No
 T2021_3.csv	95.0 MB	2023-01-18 23:21	admin	hdfs	-rw-rw-rw-		No
 T2021_4.csv	94.8 MB	2023-01-18 23:21	admin	hdfs	-rw-rw-rw-		No
 T2021_5.csv	94.8 MB	2023-01-18 23:22	admin	hdfs	-rw-rw-rw-		No
 T2021_6.csv	95.2 MB	2023-01-18 23:22	admin	hdfs	-rw-rw-rw-		No
 T2021_7.csv	94.8 MB	2023-01-18 23:22	admin	hdfs	-rw-rw-rw-		No
 T2021_8.csv	95.2 MB	2023-01-18 23:23	admin	hdfs	-rw-rw-rw-		No
 T2021_9.csv	94.7 MB	2023-01-18 23:23	admin	hdfs	-rw-rw-rw-		No
 T2022_1.csv	95.0 MB	2023-01-19 00:01	admin	hdfs	-rw-rw-rw-		No
 T2022_10.csv	95.1 MB	2023-01-19 00:04	admin	hdfs	-rw-rw-rw-		No
 T2022_11.csv	74.7 MB	2023-01-19 00:04	admin	hdfs	-rw-rw-rw-		No
 T2022_2.csv	94.9 MB	2023-01-19 00:01	admin	hdfs	-rw-rw-rw-		No
 T2022_3.csv	94.9 MB	2023-01-19 00:01	admin	hdfs	-rw-rw-rw-		No
 T2022_4.csv	94.8 MB	2023-01-19 00:02	admin	hdfs	-rw-rw-rw-		No

Dashboard

Services

HDFS

YARN

MapReduce2

Tez

Hive

HBase

Pig

Sqoop

Oozie

ZooKeeper

Storm

Infra Solr

Atlas

Kafka

Knox

Name	Size	Last Modified	Owner	Group	Permission	Erasure Coding	Encrypted
T2021_8.csv	95.2 MB	2023-01-18 23:23	admin	hdfs	-rw-rw-rw-		No
T2021_9.csv	94.7 MB	2023-01-18 23:23	admin	hdfs	-rw-rw-rw-		No
T2022_1.csv	95.0 MB	2023-01-19 00:01	admin	hdfs	-rw-rw-rw-		No
T2022_10.csv	95.1 MB	2023-01-19 00:04	admin	hdfs	-rw-rw-rw-		No
T2022_11.csv	74.7 MB	2023-01-19 00:04	admin	hdfs	-rw-rw-rw-		No
T2022_2.csv	94.9 MB	2023-01-19 00:01	admin	hdfs	-rw-rw-rw-		No
T2022_3.csv	94.9 MB	2023-01-19 00:01	admin	hdfs	-rw-rw-rw-		No
T2022_4.csv	94.8 MB	2023-01-19 00:02	admin	hdfs	-rw-rw-rw-		No
T2022_5.csv	94.9 MB	2023-01-19 00:03	admin	hdfs	-rw-rw-rw-		No
T2022_6.csv	94.8 MB	2023-01-19 00:03	admin	hdfs	-rw-rw-rw-		No
T2022_7.csv	95.2 MB	2023-01-19 00:03	admin	hdfs	-rw-rw-rw-		No
T2022_8.csv	94.8 MB	2023-01-19 00:03	admin	hdfs	-rw-rw-rw-		No
T2022_9.csv	95.0 MB	2023-01-19 00:04	admin	hdfs	-rw-rw-rw-		No
articles.csv	34.5 MB	2023-01-22 00:39	admin	hdfs	-rw-r--r--		No
customers.csv	197.5 MB	2023-01-22 00:40	admin	hdfs	-rw-r--r--		No

No

Untitled Note 1

▶ ⌵ 📖 ✎ 📄 ⬇

Q 🗑

🖨 ⚙ 🔒 default ▾

FINISHED ▶ ⌵ 📖 ⚙

```
%spark2
val hive = new org.apache.spark.sql.Session.Builder().getOrCreate()
val data = spark.read.format("csv").option("header", "true").load("hdfs:///tmp/sample/articles.csv")
data.createOrReplaceTempView("articles")
hive.sql("select * from articles limit 10").show()
```

SPARK JOBS FINISHED ▶ ⌵ 📖 ⚙

article_id	product_code	prod_name	product_type_no	product_type_name	product_group_name	graphical_appearance_no	graphical_appearance_name	colour_group_code	colour_group_name	perceived_colour_value_id	perceived_colour_value_name	perceived_colour_master_id	perceived_colour_master_name	department_no	department_name	index_code	index_name	index_group_no	index_group_name	section_no	section_name	garment_group_no	garment_group_name	detail_desc
0108775015	0108775	Strap top	253	Vest top	Garment Upper body	1010016	Solid	09	Black	4	Dark	5	Black	1676	Jersey Basic	A	Ladieswear	1	Ladieswear	16	Womens Everyday B...	1002	Jersey Basic	
0108775044	0108775	Strap top	253	Vest top	Garment Upper body	1010016	Solid	10	White	3	Light	9	White	1676	Jersey Basic	A	Ladieswear	1	Ladieswear	16	Womens Everyday B...	1002	Jersey Basi	
0108775051	0108775	Strap top (1)	253	Vest top	Garment Upper body	1010017	Stripe	11	Off White	1	Dusty Light	9	White	1676	Jersey Basic	A	Ladieswear	1	Ladieswear	16	Womens Everyday B...	1002	Jerse	

Took 15 sec. Last updated by anonymous at January 21 2023, 10:45:39 PM.

```
%sql
SELECT * FROM articles
```

SPARK JOB FINISHED ▶ ⌵ 📖 ⚙